

Integrating Scientific Simulations with Machine Learning Algorithms

Ludger Paehler
Jan Hückelheim
Sri Hari Krishna Narayanan

SIAM CSE 2023



U.S. DEPARTMENT OF
ENERGY

Office of
Science

Welcome



Ludger
Paehler



Jan
Hückelheim



Krishna
Narayanan

► Thanks for showing up on the last day!

What do we have in store

- ▶ Session 1:
 - ▶ Introduction
 - ▶ Seed matrices
 - ▶ Demo & Hands on: AD basics in PyTorch, Tapenade & Enzyme
- ▶ Session 2:
 - ▶ Interfacing in PyTorch
 - ▶ Demo & Hands on: Example 1
 - ▶ Demo & Hands on: Example 2

Resources: <https://tinyurl.com/siamcse23>

Our goals for this tutorial

Participants will . . .

- ▶ be able to use AD tools effectively,
- ▶ imagine what ML frameworks do “behind the scenes”,
- ▶ understand how ML frameworks and simulation code can be interfaced
- ▶ understand enough about AD concepts that you can diagnose problems


What we assume about you

Computational scientists facing challenging problems, through advanced modeling and simulation, using the most capable computers.

- ▶ large complex codes
- ▶ continuously developed
- ▶ including sophisticated math/physics
- ▶ using multiple libraries
- ▶ performance is essential
- ▶ parallelism involved
- ▶ familiar with ML frameworks

Why derivatives?

Sensitivity Analysis:

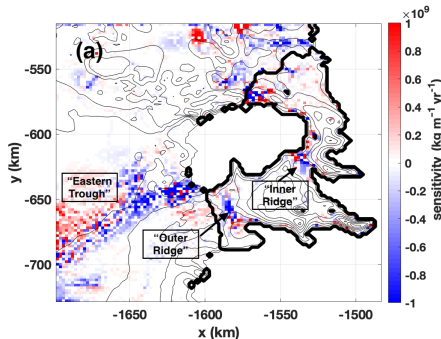


Logarithms of absolute value of adjoint sensitivities, for the Antarctic Ice Sheet. Control variables are the [A] initial ice thickness, [B] mean January precipitation [C] surface temperature, and [D] basal temperature.

Find sensitivity of the computed field *wrt* one input parameter

Why derivatives?

Sensitivity Analysis:




Sensitivity of total (area-integrated) melt to bathymetry in Dotson-Crosson experiment.

Find sensitivity of the computed field *wrt* one input parameter

Why derivatives?

Sensitivity Analysis:




Adjoint sensitivity of loss of VAF to basal melting under the ice shelves adjacent to Smith Glacier.


Find sensitivity of the computed field *wrt* one input parameter

Why derivatives?


Gradient-driven optimization



before:



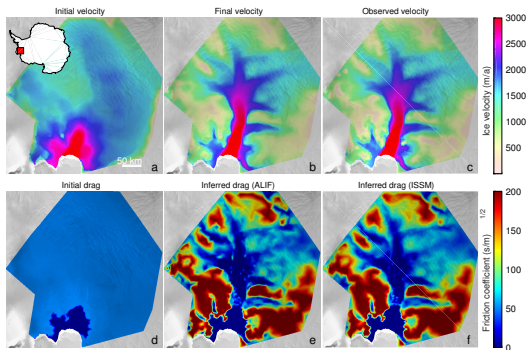
after:



Why derivatives?

Inverse problems:

from measurements and model, estimate hidden parameters



Other uses: Reduced models, Error estimation, Mesh adaption, Uncertainty Quantification, Backpropagation for ML training. . .

What is AD?

The chain rule applied to algorithms

See any algorithm/program $P: \{I_1; I_2; \dots I_p; \}$ as:

$$F: \mathbb{R}^n \rightarrow \mathbb{R}^m \quad F = f_p \circ f_{p-1} \circ \dots \circ f_1$$

Define for short:

$$V_0 = \mathbf{input} \quad \text{and} \quad V_k = f_k(V_{k-1})$$

Apply the chain rule:

$$F'(V_0) = f'_p(V_{p-1}) \times f'_{p-1}(V_{p-2}) \times \dots \times f'_1(V_0)$$

...and transform P to make it compute that.

Cost considerations

$$F'(V_0) = f'_p(V_{p-1}) \times f'_{p-1}(V_{p-2}) \times \cdots \times f'_1(V_0)$$

is often expensive:

- ▶ in computation time
- ▶ in storage space

What can save us:

- ▶ The shapes of the f'_k matter
- ▶ The final usage may not require the full F' but only a projection

Cost considerations

[illegible]

- ▶ How should we group the various matrices?
- ▶ The order matters!
- ▶ Classical projections of F'
 - ▶ $F' \times \dot{V}_0$, “forward” or “tangent” mode
 - ▶ $\overline{V_p} \times F'$ “reverse” or “adjoint” mode

Demo (with Tapenade)

```
SUBROUTINE F00(x, y)
  IMPLICIT NONE
  REAL :: x, y
  INTRINSIC SIN
  y = SIN(x)**2
END SUBROUTINE F00
```

- Can you compute the derivatives yourself?

Demo with Tapenade: tangent

```
SUBROUTINE F00_D(x, xd, y, yd)
  IMPLICIT NONE
  REAL :: x, y
  REAL :: xd, yd
  INTRINSIC SIN
  REAL :: temp
  temp = SIN(x)
  yd = 2*temp*COS(x)*xd
  y = temp*temp
END SUBROUTINE F00_D
```

- Note the common sub-expression elimination
- This computes the primal as well as the derivatives

Demo with Tapenade: adjoint

```
SUBROUTINE F00_B(x, xb, y, yb)
  IMPLICIT NONE
  REAL :: x, y
  REAL :: xb, yb
  INTRINSIC SIN
  xb = xb + COS(x)*2*SIN(x)*yb
  yb = 0.0
END SUBROUTINE F00_B
```

- ▶ The adjoint and the tangent are very similar
- ▶ Note that the primal code is not computed alongside the adjoint

Focus on forward mode: $F' \times \dot{V}_0$

$$F' \times \dot{V}_0 = f'_p(V_{p-1}) \times f'_{p-1}(V_{p-2}) \times \cdots \times f'_1(V_0) \times \dot{V}_0$$

\dot{V}_0 is a vector \Rightarrow compute from right to left!

$$F' \times \dot{V}_0 = (f'_p(V_{p-1}) \times (f'_{p-1}(V_{p-2}) \times \cdots \times (f'_1(V_0) \times \dot{V}_0) \cdots))$$

This corresponds to P's original order

\Rightarrow interleave derivative and primal computation.

Easy!




Focus on reverse mode: $\overline{V_p} \times F'$

$$\overline{V_p} \times F' = \overline{V_p} \times f'_p(V_{p-1}) \times f'_{p-1}(V_{p-2}) \times \cdots \times f'_1(V_0)$$

Vector now on the left \Rightarrow compute from left to right

$$\overline{V_p} \times F' = (\dots ((\overline{V_p} \times f'_p(V_{p-1})) \times f'_{p-1}(V_{p-2})) \times \cdots \times f'_1(V_0))$$

Not so easy, but worth the effort!




A **forward sweep**, and then a **backward sweep**

The derivative instructions form the **backward sweep**

“**Data-flow reversal**” to get V_k ’s in **reverse order**

Cost model

$$F : \text{in} \in \mathbb{R}^n \rightarrow \text{out} \in \mathbb{R}^m$$



- ▶ full F' cost grows like n using the forward mode
Good if $n \leq m$
- ▶ full F' cost grows like m using the reverse mode
Good if $n \gg m$ (e.g. $m = 1$ for a gradient)

End of basics

Let's look in detail !